

# phytools 2.0

An updated R ecosystem for phylogenetic comparative methods (and other things)

Liam J. Revell, University of Massachusetts Boston

26 February, 2023

## Abstract

1. Phylogenetic comparative methods comprise the general endeavor of using an estimated phylogenetic tree (or set of trees) to make secondary inferences: about trait evolution, diversification, biogeography, community ecology, and a wide variety of other phenomena or processes.
2. Over the past ten years, the *phytools* R package (Revell, 2012) has grown to become one of the most important tools for phylogenetic comparative analysis.
3. *phytools* is a diverse contributed R library now consisting of hundreds of different functions covering a wide range of methods and purposes in phylogenetic biology.
4. As of the time of writing, *phytools* included functionality for fitting models of trait evolution, for reconstructing ancestral states, for studying diversification on trees, and for visualizing phylogenies, comparative data, and fitted models, as well as for a number of other things.
5. Here, I describe *some* of the major features of and recent updates to *phytools*, but I also illustrate several popular workflows of the *phytools* computational software.

## Introduction

Phylogenetic trees are the directed graphs used to represent historical relationships among a set of operational taxa that are thought to have arisen via a process of descent with modification and branching. Operational taxa in a reconstructed phylogenetic tree might be gene copies, paralogous and orthologous members of a gene family, viral sequences, whole genomes, human cultural groups, or biological species. According to its broadest definition, the phylogenetic comparative method corresponds to the general activity of using a known or (most often) inferred phylogenetic tree to learn something *else* (apart from the relationships indicated by the tree) about the evolutionary process or past, the contemporary ecology, the biogeographic history, or the origins via diversification, of the particular operational taxa of our estimated tree (Harvey and Pagel, 1991; Harmon, 2018; Revell and Harmon, 2022).

Modern phylogenetic comparative methods are not new. Perhaps the most important article in the development of the phylogenetic approach to comparative biology (Felsenstein, 1985) was first authored nearly 40 years ago, and was even the subject of a relatively recent retrospective (Huey et al., 2019). Nonetheless, it's fair to say that phylogenetic comparative methods have seen a relatively impressive expansion and diversification over the past two decades. This has included the development of new approaches for studying the generating processes of trees (in other words, speciation and extinction), the relationship between phenotypic traits and species diversification, and a range of approaches for investigating heterogeneity in the evolutionary process over the tree of life. Phylogenetic comparative methods have also begun to be applied outside their traditional domain of evolutionary research. In particular, phylogenies and the comparative method have made recent appearances in studies on infectious disease epidemiology, virology, sociolinguistics, biological anthropology, molecular genetics, and community ecology.

The scientific computing environment R (R Core Team, 2023) is widely-used in biological research. One of the major advantages that R provides is that it empowers computational scientists and independent developers

to build functionality on top of the basic R computing platform. This functionality often takes the form of what are called *contributed R packages*: libraries of related functions built by individuals or research collaboratives not part of the core R development team. The growth of importance of R in phylogenetic biology stems entirely from contributed R package. Among these, the most important core function libraries are *ape* (Paradis et al., 2004; Popescu et al., 2012; Paradis and Schliep, 2019), *geiger* (Harmon et al., 2008; Pennell et al., 2014), *phangorn* (Schliep, 2011), and my package, *phytools* (Revell, 2012).

*phytools* is an R function library dedicated *primarily* to phylogenetic comparative analysis, but also including approaches and methodologies in a range of other domains of phylogenetic biology: not restricted to, but especially, visualization. The original article describing *phytools* is now more than 10 years old, and though I recently published a more comprehensive book on the subject of phylogenetic comparative methods in the R environment (Revell and Harmon, 2022), I nonetheless felt that it was time to provide a briefer (although this article is by no means brief) update of *phytools* (in particular) for the primary scientific literature. *phytools* has now grown to be *very large* – consisting of hundreds of functions, a documentation manual that’s over 200 pages in length, and (literally) tens of thousands of lines of computer code. As such, I thought it would be most useful to rather briefly summarize some of the functionality of the *phytools* R package in a few different areas, and then provide a limited number of example workflows for the “2.0” version of the *phytools* package.

## Overview

The *phytools* R package contains functionality in a wide diversity of different research areas of phylogenetics and phylogenetic biology. Rather than try to comprehensive survey this functionality, what I’ve elected to do instead, is briefly review a smaller number of methodological areas, and then illustrate each of these with one to three example analysis workflows, including the corresponding R code that can be used to reproduce the analysis and results presented here.

## Installing and loading *phytools*

This article assumes that most readers already have some familiarity with the R computing environment, and have previously installed contributed R packages. Nonetheless, to get started using *phytools* for the first time, the easiest way to install the package locally is by using the R *base* function called `install.packages`, which will download and install *phytools* from its CRAN page. (CRAN is an acronym for Comprehensive R Archive Network: a network of mirror repositories used both to archive and distribute R and contributed R packages.) This can be done as follows.

```
install.packages("phytools")
```

Readers undertaking phylogenetic analysis in the R environment for the first time will note that several other R packages are installed automatically in addition to the one (*phytools*) that we’d requested. These are packages upon which *phytools depends* – meaning that *phytools* uses one or multiple functions exported by each of these packages in its own R code. More will be said later about the dependency relationship between *phytools* and other packages of the R and R phylogenetic ecosystems.

Having installed *phytools*, to proceed and use it in an interactive R session, we normally load it. (Loading an R package simply makes the names of the functions visible and available in our current R session.)

```
library(phytools)
```

```
## Loading required package: ape
```

```
## Loading required package: maps
```

## Discrete characters

The *phytools* package now contains a wide range of different methods and models for the analysis of *discrete character evolution* on trees. In this section, I'll discuss just a few of these.

### Stochastic character mapping

Perhaps the most important and widely-used discrete character method of *phytools* is a method called 'stochastic character mapping' (Huelsenbeck et al., 2003). Stochastic character mapping is an analytical technique in which we randomly sample discrete character histories ("stochastic maps") of our trait on the tree under a specified model. One example of a stochastic character mapping analysis would be to first fit (e.g., using Maximum Likelihood) the character transition model (a variant of the *Mk* discrete character evolution model of Lewis, 2001; also see Harmon, 2018), and then proceed to randomly sample a set of 1,000 (perhaps) discrete character histories based on this character. (Other workflows are also popular and possible to undertake within R. For instance, rather than use a model of character evolution that has been optimized using Maximum Likelihood, one can instead sample parameters of the evolutionary process from their joint posterior probability distribution using Bayesian MCMC.)

To illustrate this workflow, I'll use a discrete, ecological trait for a small phylogeny of centrarchid fishes from Revell and Collar (2009). Since the trait (which we'll refer to as "feeding mode") is binary, meaning it only takes two levels, there are a total of four possible discrete character (extended *Mk*) models: equal back-and-forth transitions between the two character values; different rates; and then the two different irreversible trait evolution models. *phytools* now allows us to fit each of these four different models, compare them, and then pass the model weights and fitted models directly to our stochastic mapping function. Our mapping function (called `simmap`) will then proceed to automatically sample stochastic character histories with probabilities that are proportional to each model weight.

For this example, our data have been packaged with the *phytools* library. so we can load them in an interactive R session using the function `data` as follows.

```
data(sunfish.tree)
data(sunfish.data)
```

For our *Mk* model-fitting function (which here will be the *phytools* function `fitMk`), and for many other character methods of the *phytools* R package, our input data typically takes the form of a character or factor vector. Personally, I prefer to use factors, because in that case we can easily access the levels assumed by the character through a call of the *base* R function `levels`. (In this example our input data consists of a data frame in which the `feeding.mode` column is coded as a factor. In general, however, had we read this same data from an input text file in, for example, comma-separated-value format, R would have created a character, rather than factor, formatted column by default. To adjust this we can set the argument `stringsAsFactors=TRUE` in our file-reading function, which, in that case, might be the *base* R function `read.csv`.)

```
feeding_mode<-setNames(sunfish.data$feeding.mode,
  rownames(sunfish.data))
levels(feeding_mode)
```

```
## [1] "non" "pisc"
```

Here we see that our factor vector has two levels: "non" and "pisc". These two character levels refer to non-piscivorous and piscivorous fish species in this group. Since R factors have no particular character limit on their levels, let's update our data according: once again using the function `levels`. (`levels` is an odd R function in that it can serve both as an *extractor* function, that pulls out the levels of a factor; as well as acting as an assignment or *replacement* function, in which the levels of the factor are updated. When we adjust our factor levels for `feeding_mode`, we're using `levels` in this latter fashion.)

```
levels(feeding_mode)<-c("non-piscivorous", "piscivorous")
levels(feeding_mode)
```

```
## [1] "non-piscivorous" "piscivorous"
```

Now we're ready to proceed and fit our models. To do so, in this case, I'll use the *phytools* function `fitMk` and fit a total of four models, as previously indicated: "ER", the equal rates model; "ARD", the all-rates-different model in which we allow different forward and backward rates for our character; and, lastly, the two different irreversible models – one in which non-piscivory can evolve to piscivory, but not the reverse; and a second in which the opposite is true. For our irreversible models, we'll tell `fitMk` how to fit the model by creating and supplying a *design matrix* for each model. This design matrix should be of dimensionality  $k \times k$ , for  $k$  levels of the trait, with integer values in the positions of the matrix corresponding to allowed transitions and zeros otherwise. (We use different non-zero integer value for each rate that we want to permit to assume a different value.)

Since our  $k = 2$ , this is very easy; however, the same principle would apply to any value of  $k$  (Revell and Harmon, 2022).

```
ER.model<-fitMk(sunfish.tree,feeding_mode,
  model="ER")
ARD.model<-fitMk(sunfish.tree,feeding_mode,
  model="ARD")
Irr1.model<-fitMk(sunfish.tree,feeding_mode,
  model=matrix(c(0,1,0,0),2,2,byrow=TRUE))
Irr2.model<-fitMk(sunfish.tree,feeding_mode,
  model=matrix(c(0,0,1,0),2,2,byrow=TRUE))
```

Having fit our four models, we can also compare them to see which is best-supported by our data. I'll use a generic `anova` function call it pass it each of our fitted models as separate arguments in the function. `anova` will print the results of our model comparison; however, it's very important that we assign the *result* of `anova` to a new object (in my example, I'll call it `sunfish.aov`, but this is arbitrary) as follows.

```
sunfish.aov<-anova(ER.model,Irr1.model,Irr2.model,
  ARD.model)
```

```
##           log(L) d.f.      AIC    weight
## ER.model   -13.07453    1 28.14906 0.3486479
## Irr1.model -12.98820    1 27.97640 0.3800846
## Irr2.model -14.20032    1 30.40064 0.1130998
## ARD.model  -12.86494    2 29.72987 0.1581677
```

This table shows each of our fitted model names, their log-likelihoods, the number of parameters estimated for each model, a value of the Akaike Information Criterion (AIC), and the Akaike model weights. Smaller values of AIC indicate better support for the corresponding model, taking into account its parameter complexity. Model weights indicate the “weight of evidence” favoring each of our prior hypotheses.

With the result of our `anova` call in hand (as the `sunfish.aov` object), we're ready to pass it directly to the generic `simmap` method. By design, doing so will tell `simmap` to generate stochastic character maps under each of our four models with relative frequencies that are equal to the weight of evidence supporting of each model. (If we preferred, we could've just generated stochastic character maps for the best-supported of our four models. Using the `simmap` generic method, this would be done by setting the argument `weighted=FALSE`, or simply by passing our favored Mk model directly to the function.)

```
sunfish.simmap<-simmap(sunfish.aov,nsim=1000)
sunfish.simmap
```

```
## 1000 phylogenetic trees with mapped discrete characters
```

In spite of the significant number of stochastic simulations involved, this analysis should run fairly quickly (obviously, depending on the speed of our computer). In part this is because it *does not* involve estimating the Mk transition matrix, **Q**, for each sampled model. An additional significant advantage of this approach

is that it has also permitted us to (partly) account for uncertainty in our model estimation that's due to uncertainty in model selection.

Figure 1 shows a set of six, randomly chosen stochastic character histories for our trait (feeding mode) on our input phylogeny. Readers should see that each of these are consistent with our observed value of the binary trait at the tips of the tree, but each different one from the other in the specific hypothesis of trait evolution that it represents.

To create my color palette for plotting I used a different contributed R package that we haven't seen yet called *viridisLite* by Garnier et al. (2022). To replicate this plot exactly, users should first install *viridisLite* from CRAN by running `install.packages("viridisLite")`, but do not need to load it.

```
cols<-setNames(viridisLite::viridis(n=2),
  levels(feeding_mode))
par(mfrow=c(2,3))
plot(sample(sunfish.simmap,6),ftype="i",fsize=0.6,
  colors=cols,offset=0.2)
```




Figure 1: Six randomly chosen stochastic character maps of feeding mode (non-piscivorous vs. piscivorous) on a phylogeny of 28 centrarchid fish species. Stochastic character mapping involves randomly sampling character histories that are consistent with our tip data in proportion to their probability under a model. In this case, histories were sampled under the set of four alternative Mk models of a binary trait, in proportion to the weight of evidence supporting each model.

Although this already gives a sense of the uncertainty of our ancestral character history on the tree for our trait, most commonly we don't simply graph a subset (or all) of our stochastic mapped trees. Typically, instead, we first *summarize* our stochastic character maps (in multiple ways) before proceeding to plot or

analyze these summarized results.

Perhaps most often, *phytools* users undertaking stochastic character mapping want to compute the posterior probabilities of each value of the character trait at each internal node of the tree. These values are calculated using the generic `summary` method for our object class, which can then be plotted using a generic `plot` function call as follows.

```
plot(summary(sunfish.simmap), ftype="i", fsize=0.7,
     colors=cols, cex=c(0.6, 0.3))
legend("topleft", levels(feeding_mode), pch=21,
     pt.cex=1.5, pt.bg=cols, bty="n", cex=0.8)
```




Figure 2: Posterior probabilities at each ancestral node of the centrarchid tree of Figure 1 from stochastic character mapping using model weights to sample across four different extended  $Mk$  trait evolution models.

A correct interpretation of the graph of Figure 2 shows the observed discrete character states (at the tips of the tree) and the posterior probabilities from stochastic mapping that each internal node is in each state, in which transition models were sampled in proportion to the weight of evidence in support of each model.

In addition to node probabilities, *phytools* users undertaking a stochastic character mapping analysis are often interested in the number of changes of each type that is implied by the evolutionary process and our data. Stochastic mapping samples character histories and thus can be used to produce an estimate of the posterior probability distribution of the *number* of character changes of each type on the tree, under our sampled model or models.

To obtain this, we'll first call the generic function `density` which, when given an object from stochastic mapping, computes the relative frequency distribution of changes of each type. We can then graph these distributions (remember, our character is binary, so there are only two types of character state change: non-piscivorous  $\rightarrow$  piscivorous, and the reverse) using a different generic `plot` method.

To recreate this analysis completely, readers will also have to install an additional contributed R package called *coda* (Plummer et al., 2006) in the same way that we installed *phytools* earlier. If we have *coda* installed, it will be loaded automatically.

```

sunfish.density<-density(sunfish.simmap)

## Loading required package: coda
par(mfrow=c(1,2),las=1,cex.axis=0.7,cex.lab=0.8)
COLS<-setNames(cols,sunfish.density$trans)
plot(sunfish.density,ylim=c(0,0.6),transition=names(COLS)[1],
     colors=COLS[1],main="")
mtext("a) transitions to piscivory",line=1,
      adj=0,cex=0.8)
plot(sunfish.density,ylim=c(0,0.6),transition=names(COLS)[2],
     colors=COLS[2],main="")
mtext("b) transitions to non-piscivory",line=1,
      adj=0,cex=0.8)

```

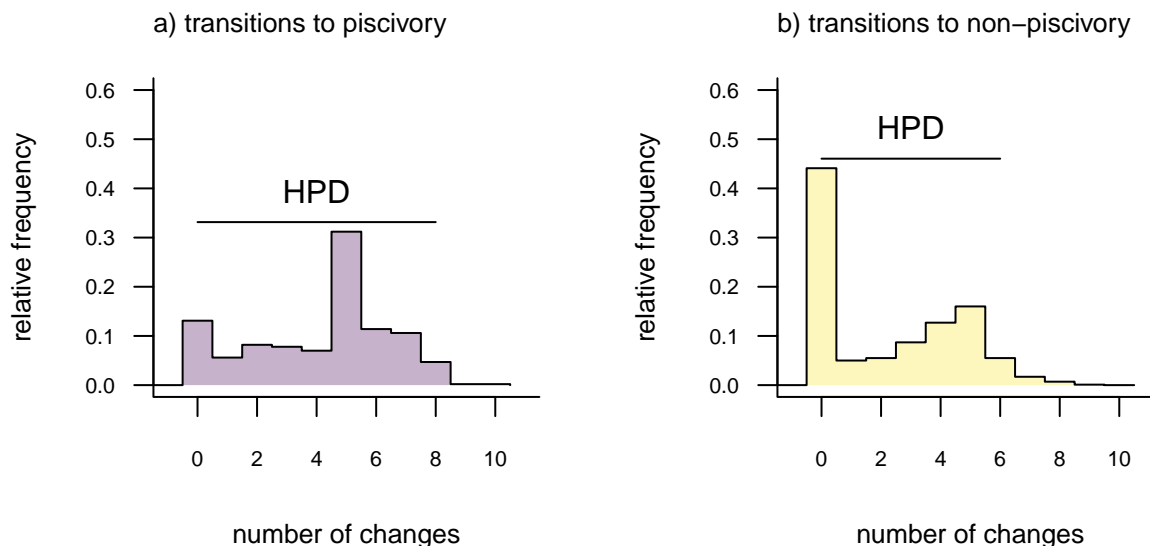


Figure 3: Posterior probability distributions of changes either (a) from non-piscivory to piscivory, or (b) from piscivory to non-piscivory, obtained from an analysis of stochastic mapping. See main text for additional details.

The distributions of Figure 3 show the relative frequencies of changes of each type across our set of mapped histories, as well as Bayesian 95% high probability density intervals from *coda*. An interesting attribute of these posterior distributions is that they are both *bi-modal*. This is due in part to our procedure of model-averaging in which we sampled both reversible and *irreversible* character evolution models in proportion to their weights. (The weight of evidence was highly similar between our equal-rates model and the irreversible model in which piscivory is acquired from piscivory, but never the reverse.)

Lastly, in addition to these analyses, *phytools* also makes it easy to visualize the posterior probabilities of each of the two trait conditions not only at nodes, but also along the branches of the phylogeny. This can be done using the function `densityMap` (Revell, 2013), which creates a graph showing the probability *density* of stochastic history in each of our mapped states. By design, in *phytools* this object can be first created, using the `densityMap` function, updated, using the function `setMap` to adjust the color palette for plotting, and then graphed, using a generic plot method designed for the object class.

```

sunfish.densityMap<-densityMap(sunfish.simap,
  plot=FALSE,res=1000)
sunfish.densityMap<-setMap(sunfish.densityMap,
  viridisLite::viridis(n=10))
plot(sunfish.densityMap,lwd=3,outline=TRUE,fsize=0.8,
  legend=0.1)

```

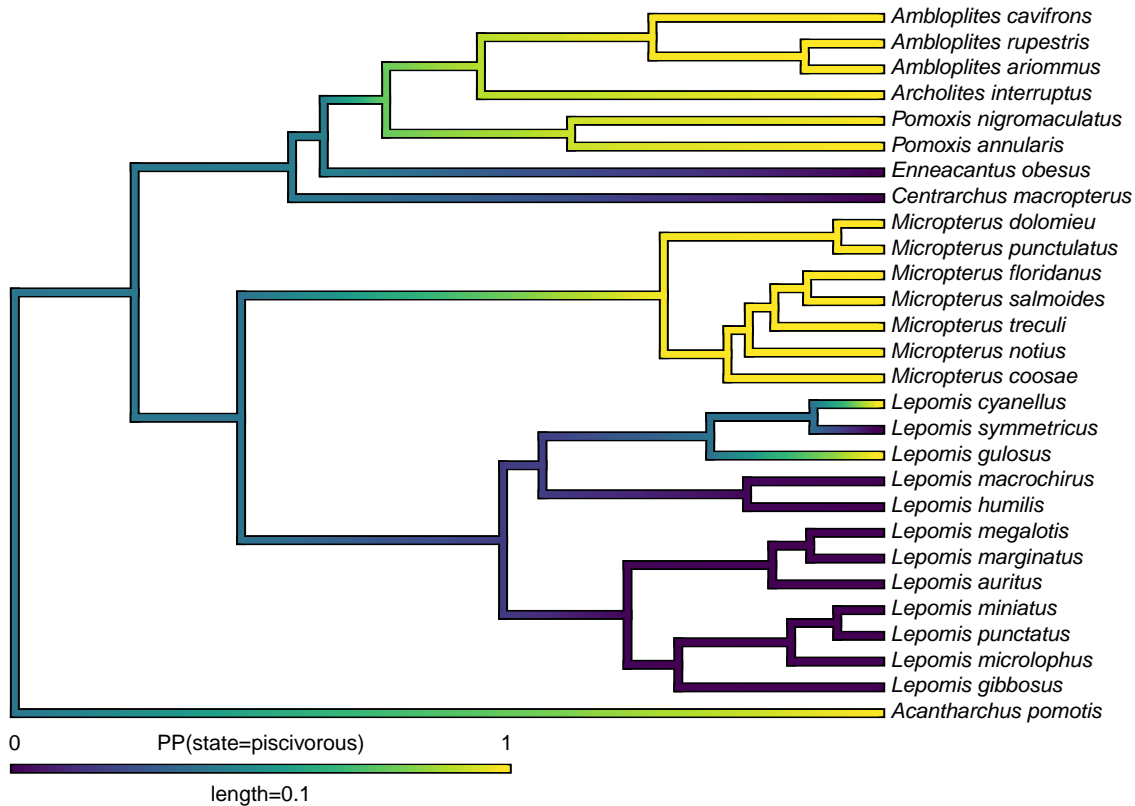


Figure 4: Posterior probability density of each of the two character levels, piscivory and non-piscivory, mapped along the edges of a tree of centrarchid fishes. See main text for more details.

## The polymorphic trait evolution model

Another important, but much more recently-added, tool in the *phytools* R package is a method (denominated `fitpolyMk`) designed to fit a discrete character evolution model to trait data containing intraspecific polymorphism. In this case, the model is one in which an evolutionary transition from (say) character state  $a$  to character state  $b$  must first pass through the intermediate polymorphic condition of  $a + b$ . This model becomes increasingly complicated for multi-state data in which we must also consider whether our characters are evolving in an ordered or unordered fashion. Figure 5 shows the general structure of an *ordered* polymorphic trait evolution model (in panel a) and an *unordered* model (in panel b), both for the same number of monomorphic conditions of our trait.

```

par(mfrow=c(1,2))
graph.polyMk(k=4,ordered=TRUE,states=0:3,mar=rep(0.1,4))
mtext("a) ordered polymorphic model",line=-1,adj=0.2,cex=0.8)
graph.polyMk(k=4,ordered=FALSE,states=letters[1:4],

```



```
mar=rep(0.1,4),spacer=0.15)
mtext("b) unordered polymorphic model",line=-1,adj=0.2,cex=0.8)
```




Figure 5: Example structures of two alternative polymorphic trait evolution models for characters with four monomorphic conditions: a) an ordered model; b) an unordered model. See main text for additional details.

Obviously, the potential parameter complexity of the unordered polymorphic trait evolution model is *higher* than the ordered model. Since the unordered model, has all ordered models as a special case, ordered and unordered models can also be compared using likelihood-ratio tests or information criteria.

To try out our polymorphic trait evolution model, let's use an excellent, recently-published dataset from Halali et al. (2020) consisting of a phylogenetic tree containing 287 *Mycalesina* butterfly species and data for habitat use. Habitat use was coded as a polymorphic trait in which, for example, a species using both "forest" and forest "fringe" habitat would be coded as "forest+fringe". In this case, our polymorphic trait evolution will assume that to evolve from forest specialization to fringe specialization, a species must first (at least transiently) evolve through the polymorphic state of using both habitats at once. This seems logical.

Our data are packaged with the *phytools* library, so can be loaded using the `data` function as follows.

```
data(butterfly.tree)
data(butterfly.data)
```

Let's extract our habitat use data as a vector and print the levels that it assumes.

```
butterfly.habitat<-setNames(butterfly.data$habitat,
  rownames(butterfly.data))
print(levels(butterfly.habitat))
```

```
## [1] "forest"          "forest+fringe"    "forest+fringe+open"
## [4] "fringe"          "fringe+open"      "open"
```

Let's proceed to fit our discrete character, polymorphic trait evolution model to these data. In this instance, I'll fit a grand total of six different models. (This is not a comprehensive set of the conceivable models for polymorphic data with these levels, but it seemed like a reasonable selection for illustrative purposes.)

The first three of these models suppose that the evolution of my discrete character is totally unordered. Among these, we'll imagine, first, equal transition rates between all monomorphic states or polymorphic conditions. In our second model, we'll permit all possible transition rates between states or state combinations

to assume different values. Finally, for our third model we'll assume that the acquisition of polymorphism (or its increase) occurs with one rate, whereas the loss (or decrease) of polymorphism occurs with another, separate rate. We refer to this model as the “transient model” following Revell and Harmon (2022). This comes from the general notion that if the rate of loss exceeds the rate of gain, then polymorphism will typically be *transient* in nature. Since polymorphism tends to be less frequently observed in the type of data that typifies many phylogenetic comparative study, including this model in our set seems sensible.

To get our remaining three models, for each of the three listed above in which character evolution is unordered, we'll add a second *ordered* model where we assume that character evolution for our three monomorphic conditions tends to proceed as follows: *forest*  $\leftrightarrow$  *fringe*  $\leftrightarrow$  *open*, not forgetting the intermediate polymorphic conditions in between each of these monomorphic states!

To fit these three different models in R, we can use the function `fitpolyMk` from the *phytools* package as follows.

```
butterfly.ER_unordered<-fitpolyMk(butterfly.tree,butterfly.habitat,
  model="ER")

##
## This is the design matrix of the fitted model. Does it make sense?
##
##               forest fringe open forest+fringe forest+open fringe+open
## forest                0     0   0             1             1             0
## fringe                0     0   0             1             0             1
## open                  0     0   0             0             1             1
## forest+fringe         1     1   0             0             0             0
## forest+open           1     0   1             0             0             0
## fringe+open           0     1   1             0             0             0
## forest+fringe+open    0     0   0             1             1             1
##
##               forest+fringe+open
## forest                      0
## fringe                      0
## open                        0
## forest+fringe               1
## forest+open                 1
## fringe+open                 1
## forest+fringe+open          0
```

By default, `fitpolyMk` prints out the *design matrix* of the model that it's about to fit. This is helpful, because we should find that it corresponds with the design matrix that was discussed under the simpler Mk model of the previous section, as well as with the graphed models of Figure 5.

Nonetheless, we can turn off this message by simply setting `quiet=TRUE`. Let's do that for our remaining two unordered models.

```
butterfly.ARD_unordered<-fitpolyMk(butterfly.tree,butterfly.habitat,
  model="ARD",quiet=TRUE)
butterfly.transient_unordered<-fitpolyMk(butterfly.tree,
  butterfly.habitat,model="transient",quiet=TRUE)
```

We're not done fitting models yet, but to see how we're doing so far, we can compare our three fitted models using the generic `anova` method of their object class, as follows.

```
anova(butterfly.ER_unordered,butterfly.ARD_unordered,
  butterfly.transient_unordered)

##               log(L) d.f.      AIC      weight
## butterfly.ER_unordered    -355.8122      1 713.6244 5.048749e-16
```

```
## butterfly.ARD_unordered      -303.5900    18 643.1800 1.000000e+00
## butterfly.transient_unordered -353.4496     2 710.8991 1.972328e-15
```

This tells us that among the models we've seen so far, the best-supported is our parameter-rich all-rates-different ("ARD") model.

Now we can proceed to do the same thing, but this time setting the argument `ordered=TRUE`.

Here we should also specify the order level using the argument `order`. If `order` isn't indicated, `fitpolyMk` will simply assume that our characters are ordered alphanumerically – but this is very rarely likely to be correct! (Readers may notice that it happens to be in this case. I set the argument `order` anyway!)

```
levs<-c("forest","fringe","open")
butterfly.ER_ordered<-fitpolyMk(butterfly.tree,butterfly.habitat,
  model="ER",ordered=TRUE,order=levs,quiet=TRUE)
butterfly.ARD_ordered<-fitpolyMk(butterfly.tree,butterfly.habitat,
  model="ARD",ordered=TRUE,order=levs,quiet=TRUE)
butterfly.transient_ordered<-fitpolyMk(butterfly.tree,
  butterfly.habitat,model="transient",ordered=TRUE,
  order=levs,quiet=TRUE)
```

Now, with all six models in hand, let's compare them using a second `anova` call as follows. This time I'll save my results to the object `butterfly.aov`.

```
butterfly.aov<-anova(butterfly.ER_ordered,butterfly.ER_unordered,
  butterfly.transient_ordered,butterfly.transient_unordered,
  butterfly.ARD_ordered,butterfly.ARD_unordered)
```

```
##               log(L) d.f.      AIC      weight
## butterfly.ER_ordered    -329.0390     1 660.0779 1.639410e-09
## butterfly.ER_unordered  -355.8122     1 713.6244 3.865519e-21
## butterfly.transient_ordered -329.0205     2 662.0409 6.143598e-10
## butterfly.transient_unordered -353.4496     2 710.8991 1.510091e-20
## butterfly.ARD_ordered    -297.8100    12 619.6201 9.999923e-01
## butterfly.ARD_unordered   -303.5900    18 643.1800 7.656389e-06
```

Our model comparison shows that (among the models in our set) the best supported by far is the ordered, all-rates-different model. *phytools* has a function to graph this model, so let's proceed and do so.

```
plot(butterfly.ARD_ordered,asp=0.65,mar=rep(0.1,4),
  cex.traits=0.8)
legend("bottomleft",legend=c(
  paste("log(L) =",round(logLik(butterfly.ARD_ordered),2)),
  paste("AIC =",round(AIC(butterfly.ARD_ordered),2))),bty="n",
  cex=0.8)
```

Just as with our fitted Mk models from the prior section, we can *also* pass this fitted model object our generic stochastic character mapping method, `simmap`. If we do, it will automatically generate a set of 100 stochastic character maps under our fitted model. (We could've likewise passed `simmap` our `anova` results, but in this case nearly all the weight of evidence fell on one model: our ordered, all-rates-different model.)

```
butterfly.simmap<-simmap(butterfly.ARD_ordered)
butterfly.simmap
```

```
## 100 phylogenetic trees with mapped discrete characters
```

Let's compute a summary of this set of stochastically mapped character histories.

```
butterfly.summary<-summary(butterfly.simmap)
```




Figure 6: Best-fitting polymorphic trait evolution model for the evolution of habitat use in *Mycalesina* butterflies.

Just as we saw before, the object from our generic `summary` call can be plotted pretty easily with *phytools*.

Using the *base* graphics function `rgb`, I'll try to select colors for plotting that are evenly spaced in a red-green-blue color space in which the "corners" (red, green, and blue) correspond to the three monomorphic states of our data.

```
hab.cols<-setNames(c(rgb(0,1,0),rgb(0,0.5,0.5),rgb(1/3,1/3,1/3),
  rgb(0,0,1),rgb(0.5,0.5,0),rgb(1,0,0)),levels(butterfly.habitat))
par(fg="transparent")
plot(butterfly.summary,type="fan",ftype="off",colors=hab.cols,
  cex=c(0.4,0.2),part=0.5,lwd=1)
par(fg="black")
legend("topleft",names(hab.cols),pch=21,pt.bg=hab.cols,pt.cex=1.5,
  cex=0.8,bty="n")
```

Figure 7 shows both the observed (at the tips) and reconstructed (at the internal nodes) marginal probabilities for each of our states and polymorphic conditions.

Lastly, let's graph the posterior distribution of the *accumulation* of lineages in each state over time, using the *phytools* function `ltt` as follows. We can do that using the same color palette as in Figure 7. (This type of graph is only especially meaningful for the situation in which the taxa of our phylogeny have been completely sampled. Though this is unlikely to be true here, it should be nonetheless interesting to see.)

```
butterfly.ltt<-ltt(butterfly.simap)
par(mar=c(5.1,4.1,1.1,1.1))
plot(butterfly.ltt,show.total=FALSE,bty="n",las=1,cex.axis=0.7,
  cex.lab=0.8,colors=hab.cols)
```

The plot of Figure @ref(fig;ltt-fitpolyMk) simultaneously shows the accumulation of lineages in each mono- or polymorphic state, but also the variation attributable to uncertainty in the evolutionary history of our group.




Figure 7: Posterior probabilities of monomorphic or polymorphic conditions at internal nodes from stochastic mapping under an ordered, ARD model of trait evolution. See main text for additional details.




Figure 8: Lineage-through-time plot showing the reconstructed accumulation of lineages in each polymorphic condition or monomorphic state over time, from 100 stochastic character maps.

## Continuous characters

Numerous continuous trait methods exist in *phytools*. For better or for worse, perhaps the most popular among these is the (arguably somewhat dubious) measurement of phylogenetic signal (Revell et al., 2008). In addition to phylogenetic signal, *phytools* contains methodology for multivariate and multi-rate Brownian motion evolution, and for ancestral state reconstruction. Here I'll start by illustrating the measurement of phylogenetic signal, I'll demonstrate Bayesian ancestral state estimation, I'll show how to fit a variable-correlation multivariate Brownian trait evolution model, and, finally, I'll show a relatively new multi-rate trait evolution model that uses the estimation technique of *penalized likelihood*. This is by no means comprehensive of the range of continuous character methods implemented in the *phytools* R package; however, it nonetheless provides a relatively broad sample.

### Phylogenetic signal

Maybe the *simplest* phylogenetic comparative analysis that we might choose to undertake for a continuous trait in R is the measurement of phylogenetic signal (Revell et al., 2008).

Phylogenetic signal has been defined in various ways, but could be considered to be the simple tendency of more closely related species to bear more similarity (one to another) than they do to more distant taxa. Phylogenetic signal can be measured in multiple ways but undoubtedly the two most popular metrics are Blomberg et al.'s (2003)  $K$  statistic, and Pagel's (1999)  $\lambda$ . Conveniently, both can be calculated using the *phytools* package.

To get started in our undertaking, let's load some data from *phytools* consisting of a phylogenetic tree of *Anolis* lizards and a data frame of morphological traits. Both derive from an article by Mahler et al. (2010).

```
data(anoletree)
data(anole.data)
```

Having loaded these data, we can next extract one variable from our data array.

Phylogenetic signal can be measured for any continuous trait, but we'll use overall body size: here represented by "snout-to-vent-length" (SVL), and given on a log scale.

```
anole.bodysize<-setNames(anole.data$SVL,rownames(anole.data))
```

We can next compute our  $K$  value of Blomberg et al. (2003) using the *phytools* function `phylosig`. `phylosig` calculates  $K$  by default (that is, without specifying an argument for `method`), but if I add the argument value `test=TRUE`, `phylosig` will also conduct a statistical test of the measured value of  $K$  by comparing it to a null distribution for  $K$  obtained by permuting our observed trait values randomly across the tips of the phylogeny.

```
anole.Blomberg_K<-phylosig(anoletree,anole.bodysize,
  test=TRUE)
anole.Blomberg_K
```

```
##
## Phylogenetic signal K : 1.69526
## P-value (based on 1000 randomizations) : 0.001
```

$K$  has an expected value of 1.0 under Brownian motion (Blomberg et al., 2003). Lower values would thus indicate less phylogenetic signal than expected under Brownian evolution; whereas higher values indicate more.

Our result shows us that measured phylogenetic signal of body size in *Anolis* lizards is greater than expected under Brownian motion. One occasional source of confusion here can sometimes be that a statistical significant value of  $K$  obtained using `phylosig` *does not* indicate whether or not  $K$  is significantly greater than Brownian motion! Our statistical test, remember, is conducted via comparison to a distribution obtained via random permutations of our data vector. As such, a statistically significant result merely indicates that

the phylogenetic signal of our data is higher than we'd expect to find in data that were entirely random with respect to the tree!

In addition to Blomberg et al.'s  $K$ , *phytools* also can be used to estimate Pagel's (1999)  $\lambda$  statistic.  $\lambda$  measures phylogenetic signal as a simple scalar multiplier of the correlations of related taxa in our tree. That is to say, if  $\lambda$  has a value less than 1.0, this indicates that related species in our phylogeny have a lower degree of "autocorrelation" than expected under Brownian evolution. A value of  $\lambda$  close to zero could be taken to indicate that related species are not autocorrelated at all! We use Maximum Likelihood to find the value of  $\lambda$  that makes our data most probable. Since we can also compute a likelihood for  $\lambda = 0$ , we can test the null hypothesis of *no* phylogenetic signal in our data by computing a likelihood ratio.

```
anole.Pagel_lambda<-phylosig(anoletree,anole.bodysize,
  method="lambda",test=TRUE)
anole.Pagel_lambda
```

```
##
## Phylogenetic signal lambda : 0.999934
## logL(lambda) : 5.25344
## LR(lambda=0) : 97.8981
## P-value (based on LR test) : 4.40483e-23
```

Unsurprisingly, this tells us that we've found significant phylogenetic signal in our trait by both measures!

Along with the simple calculation of phylogenetic signal, *phytools* also contains several methods to visualize our results. In particular, for Blomberg et al.'s  $K$  we can plot the permutation distribution of  $K$  as well as our observed measure. For Pagel's  $\lambda$  we can plot the likelihood surface, our Maximum Likelihood solution, and the likelihood of  $\lambda = 0$ : the null hypothesis of our statistical tests.

Both of these plots are illustrated in Figure @ref{phylosig} for our *Anolis* body size data.

```
par(mfrow=c(1,2),cex=0.9)
plot(anole.Blomberg_K,las=1,cex.axis=0.9)
mtext("a",adj=0,line=1)
plot(anole.Pagel_lambda,bty="n",las=1,
  cex.axis=0.9,xlim=c(0,1.1))
mtext("b",adj=0,line=1)
```

## Bayesian ancestral state estimation

The *phytools* package contains multiple functions for discrete and continuous character ancestral state estimation under multiple models.

Earlier, we reviewed the method of stochastic character mapping which is an important method for ancestral state reconstruction of discrete characters.

Among the variety of approaches for ancestral character estimation of continuous characters is the function `anc.Bayes`. As its name would suggest, `anc.Bayes` performs ancestral state estimation using Bayesian MCMC. As any proper Bayesian method should, the implementation of this method allows us to include prior information about the states at internal nodes. Here, I'll illustrate the simplest type of analysis we can undertake using this function in which we use the default node priors and MCMC conditions.

To that end, I'll load a *phytools* dataset consisting of a phylogeny of cordylid lizards from Broeckhoven et al. (2016).

```
data(cordylid.tree)
data(cordylid.data)
```

Our trait data are species scores for three different principal component axes from a phylogenetic principal components analysis. PC 1 in Broeckhoven et al. (2016) separate the most lightly armored cordylids (large

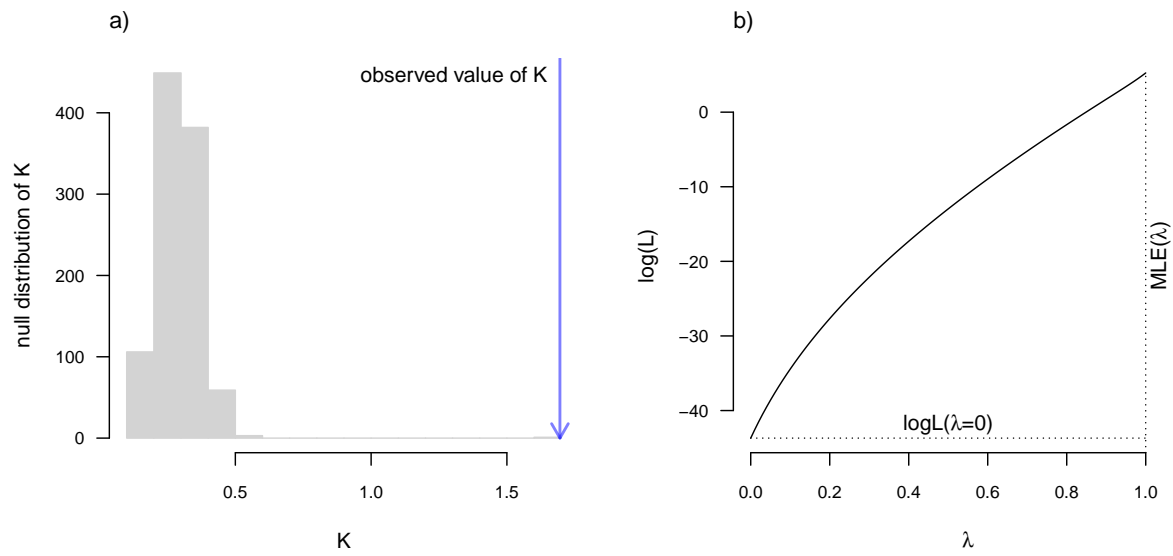


Figure 9: a) Blomberg et al. (2003) measured value of the  $K$  statistic for phylogenetic signal, compared to a null distribution of  $K$  obtained via randomization. b) Pagel's (1999)  $\lambda$  statistic for phylogenetic signal, also showing the likelihood surface.

negative)

```
cordylid.armor_score<-setNames(cordylid.data$pPC1,
                               rownames(cordylid.data))
```

```
cordylid.mcmc<-anc.Bayes(cordylid.tree,cordylid.armor_score,
                        ngen=200000)
```

```
## Control parameters (set by user or default):
```

```
## List of 7
```

```
## $ sig2 : num 0.713
```

```
## $ a : num [1, 1] 0.000422
```

```
## $ y : num [1:26] 0.000422 0.000422 0.000422 0.000422 0.000422 ...
```

```
## $ pr.mean: num [1:28] 1000 0 0 0 0 0 0 0 0 0 ...
```

```
## $ pr.var : num [1:28] 1e+06 1e+03 1e+03 1e+03 1e+03 1e+03 1e+03 1e+03 1e+03 1e+03 ...
```

```
## $ prop : num [1:28] 0.00713 0.00713 0.00713 0.00713 0.00713 0.00713 ...
```

```
## $ sample : num 100
```

```
## Starting MCMC...
```

```
## Done MCMC.
```

```
cordylid.ace<-summary(cordylid.mcmc)
```

```
##
```

```
## Object of class "anc.Bayes" consisting of a posterior
```

```
## sample from a Bayesian ancestral state analysis:
```

```
##
```

```
## Mean ancestral states from posterior distribution:
```

```
##      29      30      31      32      33      34      35      36
```

```
## -0.083829 -0.142342 -0.134643 0.080201 0.187615 0.275808 0.300722 0.369807
```



```
##          37          38          39          40          41          42          43          44
## 0.486427 0.564553 0.073103 0.044926 0.446893 0.421691 0.374172 0.246196
##          45          46          47          48          49          50          51          52
## -1.504912 -1.863289 -0.164848 -0.508119 -0.828332 -1.026121 -1.069857 0.423955
##          53          54          55
## 0.511367 -0.004952 0.014412
##
## Based on a burn-in of 40000 generations.
```

Once we obtain our Bayesian ancestral states for internal nodes, it's a straightforward thing to visualize them on the branches and nodes of the tree using the popular *phytools* plotting function `contMap` (Revell, 2013). By default, `contMap` uses Maximum Likelihood to compute ancestral states – but it can also be supplied with user-specified values, as we'll do here using our Bayesian estimates.

```
cordylid.contMap<-contMap(cordylid.tree,
  cordylid.armor_score,anc.states=cordylid.ace,
  plot=FALSE)
cordylid.contMap<-setMap(cordylid.contMap,
  viridisLite::viridis(n=10,direction=-1))
plot(cordylid.contMap,ftype="i",fsize=c(0.6,0.8),
  leg.txt="PC 1 (increasing armor)",lwd=3)
nodelabels(frame="circle",bg="white",cex=0.6)
```




Figure 10: Reconstructed ancestral values from Bayesian MCMC projected onto the nodes and edges of the tree. Numerical values at internal nodes are node indices from our input phylogeny.

(For fun, compare Figure 10 to Figure 2 of Broeckhoven et al. in which estimated ancestral state values were assigned to each branch using a similar color gradient!)

In addition to this simple analysis, we can (naturally) extract and plot posterior probability densities from *any* of our internal nodes in the tree. Let's extract and graph the posterior probability distribution of the node labeled "49" in our plot. This is the common ancestor of the *Pseudocordylus* clade, which is among the *least* heavily armored of all cordylids in our tree.

```

cordylid.node49<-density(cordylid.mcmc,what=49)
cordylid.node49

##
## Call:
## density.anc.Bayes(x = cordylid.mcmc, what = 49)
##
## Data: node 49 (1601 obs.); Bandwidth 'bw' = 0.0666
##
##      x          y
## Min.   :-2.1930  Min.   :0.000042
## 1st Qu.: -1.5523  1st Qu.:0.043395
## Median :-0.9115  Median :0.176228
## Mean   :-0.9115  Mean    :0.389782
## 3rd Qu.: -0.2708  3rd Qu.:0.736813
## Max.    : 0.3700  Max.    :1.262374

par(mar=c(5.1,4.1,1.1,2.1))
plot(cordylid.node49,las=1,bty="n",main="",
      cex.lab=0.9,cex.axis=0.8,
      xlab="PC 1 (increasing armor)",
      ylab="Posterior density",
      xlim=range(cordylid.armor_score))

```




Figure 11: Posterior probability density at node 49 of Figure 10 from Bayesian MCMC ancestral state reconstruction of PC 1 from a morphological analysis on a phylogenetic tree of cordylid lizards.

This shows our estimate of the posterior probability distribution of the ancestral node state, and should be centered on the value we projected onto the tree in Figure 10.

## Multivariate trait evolution

Along with the univariate methods we've seen so far, *phytools* contains several multivariate trait evolution models, designed for both continuous and discrete characters.

One of these is an interesting model (described in Revell and Collar, 2009; Revell et al., 2022) in which the rates and evolutionary correlations between traits are allowed to vary as a function of a set of mapped regimes on the phylogeny. The underlying motivation of this method is to test hypotheses about phylogenetic heterogeneity in the evolutionary relationship between different traits on the tree.

To illustrate our method, I'll deploy a phylogeny and dataset of tropidurid lizard species from Revell et al. (2022).

```
data(tropidurid.tree)
data(tropidurid.data)
```

In this case, our phylogeny is *already* a tree with mapped regimes. We can see this by merely printing the model object that we loaded.

```
print(tropidurid.tree, printlen=2)

##
## Phylogenetic tree with 76 tips and 75 internal nodes.
##
## Tip labels:
##  Leiocephalus_raviceps, Leiocephalus_carinatus, ...
##
## The tree includes a mapped, 2-state discrete character
## with states:
##  n_rock, rock
##
## Rooted; includes branch lengths.
```

This tells us that our phylogenetic tree contains 76 taxa and a mapped regime with two states: "n\_rock" (non-rock dwelling) and "rock" (rock-dwelling). Since *phytools* permits mapped regimes to have arbitrary names, let's rename these levels in a more informative way.

To rename my mapped traits I'll use the *phytools* function `mergeMappedTraits`. `mergeMappedTraits`, as readers can probably guess, is designed to merge the mappings of two or more traits into one – but can also be employed to simply substitute one mapping name for another, as we'll use it here.

```
tropidurid.tree<-mergeMappedStates(tropidurid.tree,"n_rock",
  "non-rock dwelling")
tropidurid.tree<-mergeMappedStates(tropidurid.tree,"rock",
  "rock-dwelling")
```

Let's plot this updated tree. When I do this, I'm going to use the *phytools* plotting function `sigmoidPhylogram` which will plot our tree using curved ("sigmoidal") linking lines. (*phytools* contains lots of cool functions like this!)

```
cols<-setNames(c("white","black"),c("non-rock dwelling","rock-dwelling"))
sigmoidPhylogram(tropidurid.tree,direction="upwards",outline=TRUE,
  colors=cols,direction="upwards",outline=TRUE,lwd=2,
  fsize=0.4,ftype="i",offset=1)
legend("bottomright",c("non-rock dwelling","rock-dwelling"),pch=22,
  pt.bg=cols,cex=0.8,pt.cex=1.2)
```

Our phenotypic trait data in `tropidurid.data` consist of a single measure of overall body size (as trait 1), and a second metric trait measuring dorsoventral flattening.




Figure 12: Phylogenetic tree of tropidurid lizard species from Revell et al. (2022).

Our hypothesis is that these two traits should scale together in non-rock dwelling lizard species: bigger lizards also tend to be deeper. However, we would like to ask if this general relationship is perhaps *decoupled* among rock-dwelling species in which the force of selection favors *increased* flattening relative to their non-rock kin. (There are biomechanical and behavioral reasons to suspect this could be so. For more information, see Revell et al., 2007.)

To test this hypothesis, we'll use the *phytools* function `evolvcv.lite` which fits a heirarchical set of models for the evolutionary rates (of each character) and evolutionary correlations (between them).

```
tropidurid.fits<-evolvcv.lite(tropidurid.tree,tropidurid.data)
```

```
## Fitting model 1: common rates, common correlation...
## Best log(L) from model 1: 52.3056.
## Fitting model 2: different rates, common correlation...
## Best log(L) from model 2: 54.3968.
## Fitting model 3: common rates, different correlation...
## Best log(L) from model 3: 55.1105.
## Fitting model 4: no common structure...
## Best log(L) from model 4: 56.2877.
```

Having fit each of four models (in this case: `evolvcv.lite` actually includes several additional models, Revell et al., 2022), we can easily compare them with a generic `anova` function call as follows.

```
anova(tropidurid.fits)
```

```
##           log(L) d.f.      AIC    weight
## model 1 52.30560    5 -94.61119 0.09221085
## model 2 54.39681    7 -94.79362 0.10101737
## model 3 55.11048    6 -98.22096 0.56057433
## model 4 56.28765    8 -96.57530 0.24619744
```

Our model comparison quite clearly indicates that "model 3" is the best-supported for our data. Indeed, this

model is one in which the evolutionary covariance between overall body size and dorsoventral flattening is *negative* among rock-dwelling lineages – compared to the positive evolutionary covariance in non-rock species and across all other models, just as we’d predicted.

```
tropidurid.fits
```

```
## Model 1: common rates, common correlation
## R[1,1] R[1,2] R[2,2] k log(L) AIC
## fitted 0.2224 0.0154 0.0589 5 52.3056 -94.6112
##
## (R thinks it has found the ML solution for model 1.)
##
## Model 2: different rates, common correlation
## R[1,1] R[1,2] R[2,2] k log(L) AIC
## non-rock dwelling 0.2025 0.0187 0.0456 7 54.3968 -94.7936
## rock-dwelling 0.3043 0.0382 0.1263
##
## (R thinks it has found the ML solution for model 2.)
##
## Model 3: common rates, different correlation
## R[1,1] R[1,2] R[2,2] k log(L) AIC
## non-rock dwelling 0.2256 0.0394 0.0588 6 55.1105 -98.221
## rock-dwelling 0.2256 -0.0354 0.0588
##
## (R thinks it has found the ML solution for model 3.)
##
## Model 4: no common structure
## R[1,1] R[1,2] R[2,2] k log(L) AIC
## non-rock dwelling 0.2108 0.0325 0.0485 8 56.2877 -96.5753
## rock-dwelling 0.2794 -0.0564 0.101
##
## (R thinks it has found the ML solution for model 4.)
```

## Variable rate Brownian motion

Lastly, *phytools* recently added a function to fit a variable-rate Brownian model using penalized likelihood.

Under this model, the rate of phenotypic trait evolution ( $\sigma^2$ , under a Brownian evolutionary process) is permitted to vary from edge to edge in the tree to a degree controlled by our  $\lambda$  shrinkage or smoothing parameter (Revell, 2021). This method has already been used to, for example, investigate rate heterogeneity differences in body size evolution between cetaceans and plesiosaurs (Sander et al., 2021).

Here, I’ll apply it to the analysis of skull size evolution in a phylogenetic tree of primates. My data for this example (packaged with *phytools*) come from a book chapter by Kirk and Kay (2004).

```
data(primate.tree)
data(primate.data)
```

Our data frame, `primate.data`, contains a number of different variables. Let’s pull out just one, `Skull_length`, and (as we do) convert it to a log scale.

```
primate.lnSkull<-setNames(log(primate.data$Skull_length),
rownames(primate.data))
```

With this input data vector and our tree we’re ready to run our penalized likelihood method. Invariably, penalized likelihood methods require the user to specify a smoothing parameter – normally denominated  $\lambda$ .  $\lambda$  determines the weight that’s assigned to the penalty term of the fitted model, in our case a measure of how

much (or little) the evolutionary rate varies from edge to edge in the tree. A large value of  $\lambda$  will penalize high rate variation between edges and thus cause us to fit a model with relatively low rate heterogeneity. Small values of  $\lambda$  allow the rate to vary with little smoothing cost. Multiple approaches, such as cross-validation, have been recommended to help us to identify a suitable value of  $\lambda$  for our data – however, *minimally* I would recommend testing multiple values of  $\lambda$  and comparing the results. Let's try that here: first using  $\lambda = 1.0$  and then swapping it for a much smaller  $\lambda = 0.1$  and much larger  $\lambda = 10$  to see how these different values of our smoothing parameter affect our findings.

Before we continue, let's create a simple projection of our trait data onto the tree. In this case, I'll use the *phytools* function `edge.widthMap` which sizes the thickness of our plotted branches in proportion to the observed or reconstructed values of our continuous trait.

```
primate.widthMap<-edge.widthMap(primate.tree,primate.lnSkull)
plot(primate.widthMap,color=palette()[2],legend="log(skull length)")
```




Figure 13: Primate skull lengths (on a log scale) projected onto the edges and nodes of the phylogeny.

The function we'll use for this exercise, `multirateBM`, performs a computationally intensive optimization.

Setting the optional argument `parallel=TRUE` will help distribute the computational burden across multiple processors, if possible.

Let's start with  $\lambda = 1.0$ .

```
primate.mBM_1<-multirateBM(primate.tree,primate.lnSkull,lamba=1,
  parallel=TRUE)
```

```
## Beginning optimization....
## Using socket cluster with 16 nodes on host 'localhost'.
## Optimization iteration 1. Using "L-BFGS-B" (parallel) optimization method.
## Best (penalized) log-likelihood so far: -267.108
## Done optimization.
```

Now we can do the same with  $\lambda = 0.1$  and 10. This time we'll turn off printing by setting the optional argument `quiet=TRUE`.

```
primate.mBM_0.1<-multirateBM(primate.tree,primate.lnSkull,
  lambda=0.1,parallel=TRUE,quiet=TRUE)
primate.mBM_10<-multirateBM(primate.tree,primate.lnSkull,
  lambda=10,parallel=TRUE,quiet=TRUE)
```

Readers should note that the *specific values* of the penalized log likelihoods are not comparable between analyses with different values of the penalty coefficient,  $\lambda$ !

Finally, let's visualize the differences and similarities between each of our three plotted models.

```
par(mfrow=c(1,3))
plot(primate.mBM_1,ftype="off",lwd=2,mar=c(0.1,0.1,2.1,0.1))
mtext(expression(paste("a) ",lambda," = 1")),adj=0.1,line=0.5)
plot(primate.mBM_0.1,ftype="off",lwd=2,mar=c(0.1,1.1,2.1,0.1))
mtext(expression(paste("b) ",lambda," = 0.1")),adj=0.1,line=0.5)
plot(primate.mBM_10,ftype="off",lwd=2,mar=c(0.1,1.1,2.1,0.1))
mtext(expression(paste("c) ",lambda," = 10")),adj=0.1,line=0.5)
```




Figure 14: Estimated rates of log(body mass) evolution under a variable-rate Brownian evolution model for different values of the smoothing parameter,  $\lambda$ .

We can see from the plot of Figure 14 that even though the specific *range* of rate variation measured depends

strongly on  $\lambda$ , the specific pattern from clade to clade on the tree is relatively consistent. This should give us some confidence that the measured rate heterogeneity is a product of real variation in the evolutionary rate for our character.

## Diversification

In addition to the methods that we've seen so far, *phytools* also contains a handful of different techniques for investigating diversification. This is not a primary focus of the package, but these methods are popular and the *phytools* implementations relatively easy to use.

### Lineage through time plots

One of the simplest phylogenetic methods for studying diversification is to simply graph the accumulation of new lineages in our reconstructed phylogeny over time since the global root of the tree. This is called a lineage-through-time plot. If we graph the number of lineages through time in a fully sampled pure-birth (that is, constant-rate speciation but no extinction) phylogenetic tree, the accumulation curve should be *exponential* – or linear on a semi-logarithmic scale.

To see how the number of lineages through time are calculated and graphed using *phytools* let's load a phylogenetic tree of lizards from the diverse South American family Liolaemidae. This phylogeny is packaged with *phytools* but was published by Esquerre et al. (2019).

```
data(liolaemid.tree)
print(liolaemid.tree,printlen=2)

##
## Phylogenetic tree with 257 tips and 256 internal nodes.
##
## Tip labels:
##   Liolaemus_abaucan, Liolaemus_koslowskyi, ...
##
## Rooted; includes branch lengths.
```

We'll create lineage-through-time graph with *phytools* over two steps. First, we'll use the *phytools* function `ltt` to compute an object of class "ltt" containing our tree and a count of the number of lineages through time from the root of the tree to the tips.

```
liolaemid.ltt<-ltt(liolaemid.tree,plot=FALSE)
liolaemid.ltt

## Object of class "ltt" containing:
##
## (1) A phylogenetic tree with 257 tips and 256 internal
##     nodes.
##
## (2) Vectors containing the number of lineages (ltt) and
##     branching times (times) on the tree.
##
## (3) A value for Pybus & Harvey's "gamma" statistic of
##     gamma = 1.8129, p-value = 0.0698.
```

From the print-out we see that in addition to the tree and the lineages through time, our object also contains a value of (and a p-value for) Pybus and Harvey's (2000)  $\gamma$  statistic.  $\gamma$  is a numerical value used to describe the general shape of the lineage through time curve. If the curve is *straight* (on a semi-log scale), then  $\gamma$  should have a value close to zero. This is what we expect under a pure-birth (speciation only) diversification process. On the other hand, significantly *positive* or significantly *negative*  $\gamma$  means that the lineage through time graph curves upward or downward towards the present day. This could mean that the rate of diversification



has changed over time, but it could also be due to past extinction or incomplete taxon sampling (Revell and Harmon, 2022).

We can see that for our liolaemid lizards,  $\gamma$  is *positive*, though not significantly so.

```
par(mar=c(5.1,4.1,1.1,2.1))
plot(liolaemid.ltt,show.tree=TRUE,lwd=2,
     log.lineages=FALSE,log="y",bty="n",las=1,
     cex.axis=0.8,cex.lab=0.9,transparency=0.1)
```




Figure 15: Lineage through time plot for phylogeny of lizards from the South American family Liolaemidae.

In general, accounting for incomplete taxon sampling in the measurement of Pybus and Harvey’s (2000)  $\gamma$  statistic is important because missing taxa will tend to pull out LTT curve *downwards* (in other words, towards more negative values of  $\gamma$ ) as we approach the tips of the tree. Fortunately, there’s a simple way to account for this. If we know the true species richness of our clade of interest, we can simulate trees that match this richness under pure-birth, randomly prune taxa to the level of missingness in our own data, and then use the distribution of  $\gamma$  values across this set of simulated, and randomly pruned trees as our null distribution for hypothesis testing. This procedure is called the “Monte Carlo constant rates” (or MCCR) test. It’s implemented in *phytools* as the function `mccr`.

Of course, since the MCCR test accounts for randomly missing taxa from our tree, we must know or hypothesize a true species richness of our clade. In this instance, we’re not too preoccupied about the precise value for liolaemid, but the *Reptile Database* puts the species richness of Liolaemidae at 340. For illustrative purposes, let’s just go with that number!

```
liolaemid.mccr<-mccr(liolaemid.ltt,rho=Ntip(liolaemid.tree)/340,
                    nsim=1000)
liolaemid.mccr
```

```
## Object of class "mccr" consisting of:
```

```
##
## (1) A value for Pybus & Harvey's "gamma" statistic of
##     gamma = 1.8129.
##
## (2) A two-tailed p-value from the MCCR test of 0.002.
##
## (3) A simulated null-distribution of gamma from 1000
##     simulations.
```

This tells us that (accounting for missing taxa) our observed value of  $\gamma$  becomes highly significantly different from that expected under pure-birth. Let's plot our results.

```
par(mar=c(5.1,4.1,0.6,2.1))
plot(liolaemid.mccr,las=1,cex.lab=0.8,cex.axis=0.7,main="")
```




Figure 16: Lineage through time plot for phylogeny of lizards from the South American family Liolaemidae.

Figure 16 shows that our measured  $\gamma$  is strongly significantly *positive* – in other words, much larger than expected by random chance under a constant-rate pure birth speciation process.

## Modeling speciation and extinction

In addition to these analysis, *phytools* can also fit simple speciation and extinction models following Nee et al. (1994; Stadler, 2012). This is done using the function `fit.bd` which also allows us to take into account an incomplete taxonomic sampling fraction. Just as with Pybus and Harvey's  $\gamma$ , incomplete sampling also has the potentially to substantially distort our estimated speciation (normally given as  $\lambda$  – a different  $\lambda$  from before!) and extinction ( $\mu$ ). In this case, ignoring (or underestimating) the missing lineages in our tree will tend to cause us to underestimate the rate of extinction, as nearly all the information about extinction comes from the most recent parts of our tree! (See Revell and Harmon, 2022 for more details.)

Fitting a birth-death model using *phytools* is easy.

```
liolaemid.bd<-fit.bd(liolaemid.tree,
  rho=Ntip(liolaemid.tree)/340)
```

```
liolaemid.bd
```

```
##
## Fitted birth-death model:
##
## ML(b/lambda) = 0.351
## ML(d/mu) = 0.1771
## log(L) = 526.451
##
## Assumed sampling fraction (rho) = 0.7559
##
## R thinks it has converged.
```

We can also fit and compare this model to a simpler, pure-birth model: also called a ‘Yule’ model.

```
liolaemid.yule<-fit.yule(liolaemid.tree,
  rho=Ntip(liolaemid.tree)/340)
liolaemid.yule
```

```
##
## Fitted Yule model:
##
## ML(b/lambda) = 0.2499
## log(L) = 521.2295
##
## Assumed sampling fraction (rho) = 0.7559
##
## R thinks it has converged.
```

```
anova(liolaemid.yule,liolaemid.bd)
```

```
##           log(L) d.f.      AIC    weight
## liolaemid.yule 521.2295    1 -1040.459 0.0144644
## liolaemid.bd   526.4510    2 -1048.902 0.9855356
```

Overall, this result tells us that a birth-death (speciation and extinction) model is much better supported than our simpler Yule model.

Lastly, the *phytools* function `fit.bd` exports a likelihood function as part of the fitted model object it returns. As such, it’s a very straightforward thing to use that function to (for example) compute and graph the likelihood surface. We can do that using the base R *graphics* function `persp`. (But R contains lots of even fancier 3D plotting methods readers might be more interested in trying!)

```
ngrid<-40
b<-seq(0.25,0.45,length.out=ngrid)
d<-seq(0.10,0.25,length.out=ngrid)
logL<-matrix(NA,ngrid,ngrid)
for(i in 1:ngrid) for(j in 1:ngrid)
  logL[i,j]<-liolaemid.bd$lik(c(b[i],d[j]))
logL[is.nan(logL)]<-min(logL[!is.nan(logL)])
par(mar=rep(0.1,4))
persp(b,d,exp(logL),shade=0.3,
  phi=45,theta=20,
  xlab="speciation rate",
  ylab="extinction rate",
  zlab="likelihood",
  border=palette()[4],expand=0.3)
```




Figure 17: Visualization of the likelihood surface.

Our graph shows the very strong *ridge* in the likelihood surface (from low  $\lambda$  and low  $\mu$ , to high  $\lambda$  and high  $\mu$ ) that invariably tends to characterize the likelihood surfaces of birth-death models and often so befuddles estimation!

## Visualization

After phylogenetic comparative methods, *phytools* is perhaps best known for its phylogeny visualization methods. We’ve seen a number of these methods deployed already throughout this article. For example, in Figures 1, 2, 3, 4, 7, and 12 I illustrated custom *phytools* plotting methods for stochastic character mapping and the analysis of stochastic character mapped trees. Likewise, in Figures 5 and 6 I demonstrate *phytools* plotting methods for fitted discrete character evolution models. In Figures 9, 10, 13, and 14 I show a variety of custom methods for visualization continuous trait evolution. Finally, in Figure 8, 15, and 16 I illustrate several methods for visualizing diversification or the results from a diversification analysis using a phylogeny. This is a sparse sample of the variety of plotting methods for phylogenies, phylogenetic comparative data, and the results of phylogenetic analysis that are implemented in the *phytools* R package. In this final section, I’ll illustrate just a few more popular plotting methods of the package that we haven’t seen in prior sections.

### Co-phylogenetic plotting

Among the most popular plotting method of the *phytools* package is the function `cophylo`, which creates co-phylogenetic plots (often called “tanglegrams”). The purpose of tanglegrams varies widely from study to study. For instance, tanglegrams are sometimes used to visually illustrate the topological similarity between two groups that are hypothesized to co-speciate: for instance, an animal host and its parasites; or a plant and its pollinators.

Equally often, however, tanglegrams are put to different purposes. For instance, tanglegrams are frequently used to visualize the similarity of alternative phylogenetic hypothesis, to identify incongruence among gene trees, and even to compare a phylogenetic history to a cluster dendrogram based on morphology.

To illustrate the *phytools* tanglegram method, I’ll use a phylogenetics tree of bat species and another of their betacoronaviruses based on Caraballo (2022).

```
data(bat.tree)
data(betaCoV.tree)
```

Assuming that our tip labels differ between our different trees (and they do in this instance), we need more than just two phylogenies to create a tanglegram – we also need a table of associations linking the tip labels of one tree to those of the other. Using data from Caraballo (2022), our data for the two trees we’ve loaded is in the *phytools* data object `bat_virus.data`. Let’s load it.

```
data(bat_virus.data)
head(bat_virus.data)
```

```
##           Bats betaCoVs
## 1  Artibeus lituratus KT717381
## 2  Artibeus planirostris MN872692
## 3  Artibeus planirostris MN872690
## 4  Artibeus planirostris MN872691
## 5  Artibeus planirostris MN872689
## 6  Artibeus planirostris MN872688
```

Inspecting just the first part of this object reveals its general structure. We can see that it consists of two columns: one for each of our two trees. The elements of the first column should match the labels of our first tree, and those of the second column the labels of our second tree. There’s no problem at all if one or the other columns has repeating names: a host can (of course) be associated with more than one parasite, and vice versa!

Now let's run our co-phylogenetic analysis. This will create, not a plot, but a "cophylo" object in which the node rotation has been optimized to maximize the tip alignment of the two trees.

```
bat.cophylo<-cophylo(bat.tree,betaCoV.tree,
  assoc=bat_virus.data)
```

```
## Rotating nodes to optimize matching...
## Done.
```

We can print this object, as follows.

```
bat.cophylo
```

```
## Object of class "cophylo" containing:
##
## (1) 2 (possibly rotated) phylogenetic trees in an object of class "multiPhylo".
##
## (2) A table of associations between the tips of both trees.
```

To plot it, we'll use the a generic plot method for the object class.

I'll go ahead and adjust a few settings of the method to make our graph look nice, and I'll use species-specific linking lines so that we can see all the virus sequences that are associated with each bat host.

```
cols<-setNames(RColorBrewer::brewer.pal(n=7,name="Dark2"),
  bat.tree$tip.label)
par(lend=3)
plot(bat.cophylo,link.type="curved",fsize=c(0.7,0.6),
  link.lwd=2,link.lty="solid",pts=FALSE,
  link.col=make.transparent(cols[bat_virus.data[,1]],0.5))
pies<-diag(1,Ntip(bat.tree))
colnames(pies)<-rownames(pies)<-names(cols)
tiplabels.cophylo(pie=pies,
  piecol=cols[bat.cophylo$trees[[1]]$tip.label],
  which="left",cex=0.2)
```




Figure 18: Co-phylogenetic plot of bat species and their associated betacoronaviruses. Associations and GenBank accession numbers from Caraballo (2022).

In general, our plot reveals a surprisingly strong association between the topology of the bats and their viruses, something that Caraballo (2022) also reported.

## Projecting a tree onto a geographic map

*phytools* can be used to project a phylogenetic tree onto a geographic map.

To see how this is done in R, we can load two datasets that come with the *phytools* package. The first is a phylogenetic tree (`mammal.tree`) from Garland et al. (1992); the second is a corresponding geographic dataset that I obtained by sampling 1-3 records at random from the online citizen science resource *iNaturalist*.

```
data(mammal.tree)
data(mammal.geog)
```

Our data (`mammal.geog`) come in the form of a *matrix*. This is important to note because our plotting function allows more than one geographic record per species and requires that our taxon labels be supplied as row names. The most common way to read data into R (for instance, using `read.table` or `read.csv`) with create a data frame, rather than a matrix, and R data frames do not permit repeating row names! The easiest way to resolve this is to read our data in with species as a data column rather than as row names, and then using our input data to create a matrix object within R.

```
head(mammal.geog)
```

```
##           lat      long
## A._alces    61.07760 -149.82472
## A._alces    45.65505  -69.65062
## A._americana 43.60828 -105.15943
## A._americana 40.73963 -104.51432
## A._buselaphus -19.96034  15.56863
## A._cervicapra 13.00615   80.24190
```

I happen to know that there are some mismatches between our input tree and data. (This is because I didn't include geographic records for all species when I was creating `mammal.geog`, but it's a common problem with comparative data!)

To identify these, I'm going to use the function `name.check` from the *geiger* comparative methods package (Pennell et al., 2014). *geiger* is not a dependent package of *phytools*, but I suspect most readers will already have it installed. If not, it can be installed from CRAN in the typical fashion.

```
library(geiger)
chk<-name.check(mammal.tree,mammal.geog)
summary(chk)
```

```
## 2 taxa are present in the tree but not the data:
##      G._granti,
##      V._fulva
##
## To see complete list of mis-matched taxa, print object.
```

This tells us that there are two taxa in our tree that are missing from our geographic data, so we can prune with out using the *ape* function `drop.tip`.

```
mammal.pruned<-drop.tip(mammal.tree,
  chk$tree_not_data)
name.check(mammal.pruned,mammal.geog)
```

```
## [1] "OK"
```

Our next step is to build the map projection that we intend to plot. This is done using the *phytools* function `phylo.to.map`. In addition to combining our phylogenetic tree and map data, this function also performs a

series of node rotations to try to optimize the alignment of our phylogeny with the geographic coordinates of our tip data. Since node rotation is arbitrary anyway, this can be helpful to facilitate a more convenient visualization.

```
mammal.map<-phylo.to.map(mammal.pruned,mammal.geog,
  plot=FALSE,quiet=TRUE)
mammal.map
```

```
## Object of class "phylo.to.map" containing:
##
## (1) A phylogenetic tree with 47 tips and 46 internal nodes.
##
## (2) A geographic map with range:
##      -85.19N, 83.6N
##      -180W, 180W.
##
## (3) A table containing 72 geographic coordinates (may include
##      more than one set per species).
##
## If optimized, tree nodes have been rotated to maximize alignment
## with the map when the tree is plotted in a downwards direction.
```

Our plotting function for this object class allows us to specify different colors for the different species being plotted. Let's do that here by choosing random colors from a divergent color palette using the CRAN package *randomcoloR* (Ammar, 2019). (Readers lacking this package should install it from CRAN – or they can elect to use a different palette, or none at all.)

```
cols<-setNames(
  randomcoloR::distinctColorPalette(Ntip(mammal.pruned)),
  mammal.pruned$tip.label)
```

Finally, we're ready to plot our tree.

```
plot(mammal.map, fsize=0.5, ftype="i",
  colors=cols, asp=0.9, lwd=1, cex.points=c(0.4,0.8))
```

In this case, the lineages of our phylogeny have a near-global distribution. For phylogenetic trees whose terminal taxa have a narrower geographic range, including phylogeographic studies, it's possible to restrict the plotted maps to countries or regions, or even to supply a custom map!

## Projecting trees into phenotypic space

In addition to projecting a phylogenetic tree onto a geographic map (as we just saw) and projecting traits onto the edges and nodes of a plotted tree, the *phytools* package also contains multiple methods to project a tree into a space defined by our traits.

```
data(vertebrate.tree)
data(vertebrate.data)

vertebrate.ace<-exp(apply(log(vertebrate.data[,c(1,3)]),2,
  fastAnc,tree=vertebrate.tree))
vertebrate.ace
```

```
##           Mass Litter_size
## 12  25.571594   15.552422
## 13  17.834793   16.114126
## 14  16.827622   14.481309
## 15   8.801275    8.791375
```



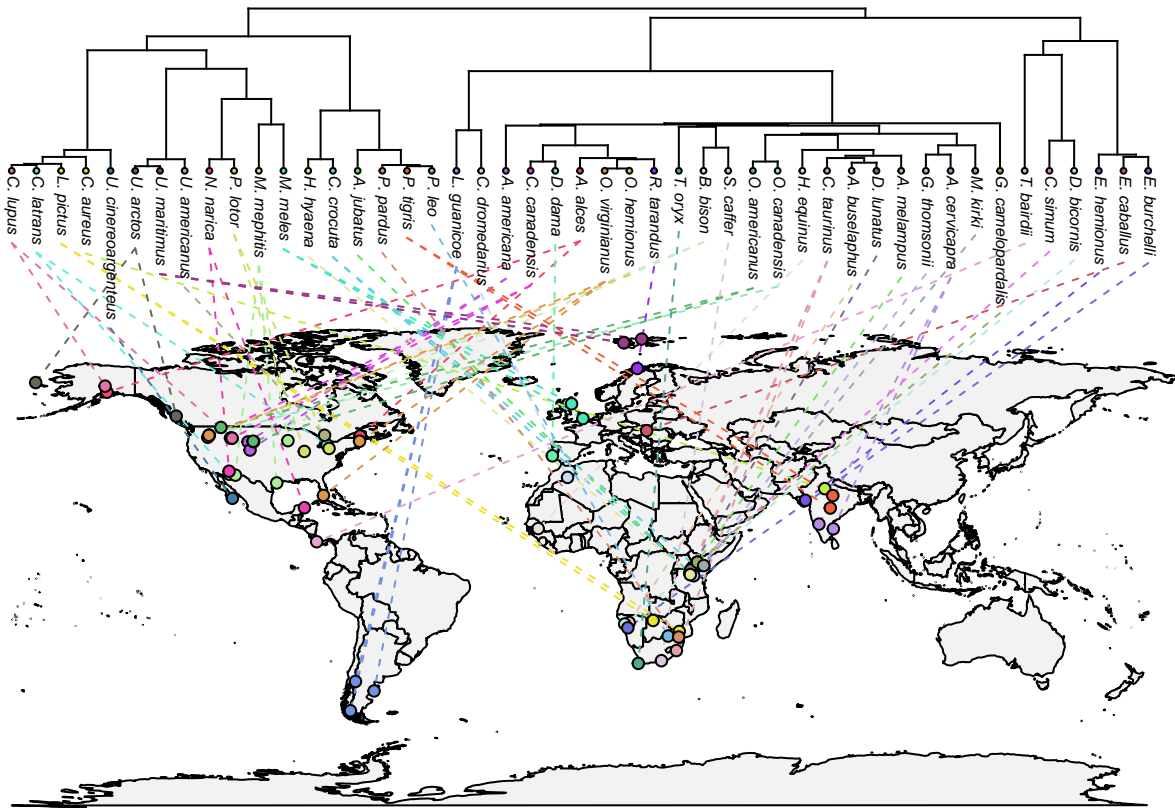


Figure 19: A phylogenetic tree of mammals projected onto a geographic map.

```
## 16 20.362215 1.653955
## 17 17.335294 1.442396
## 18 24.604666 1.610416
## 19 152.078728 1.737334
## 20 301.219076 1.582780
## 21 6.329908 9.613237

par(mar=c(5.1,4.1,0.6,2.1))
phylomorphospace(vertebrate.tree,vertebrate.data[,c(1,3)],
  A=vertebrate.ace,log="xy",xlim=c(1e-4,1e6),ylim=c(0.5,200),
  bty="n",label="off",axes=FALSE,xlab="Mass (kg)",
  ylab="Litter size",node.size=c(0,0))
axis(1,at=10^seq(-3,5,by=2),
  labels=prettyNum(10^seq(-3,5,by=2),big.mark=","),
  las=1,cex.axis=0.7)
axis(2,at=10^seq(0,2,by=1),
  labels=prettyNum(10^seq(0,2,by=1),big.mark=","),
  las=1,cex.axis=0.7)
cols<-setNames(RColorBrewer::brewer.pal(nrow(vertebrate.data),"Paired"),
  rownames(vertebrate.data))
points(vertebrate.data[,c(1,3)],pch=21,bg=cols,cex=1.2)
legend("topleft",gsub("_"," ",rownames(vertebrate.data)),
  pch=21,pt.cex=1.2,pt.bg=cols,cex=0.6,bty="n",text.font=3)
```




Figure 20: Phylomorphospace of body mass and litter size for a selection of vertebrate species.

### Plotting phenotypic data at the tips of the tree

In addition to projecting phylogenies into trait spaces, and plotting observed or reconstructed trait values on the tree, *phytools* can also help us undertake the relatively simple task of visualizing comparative trait data for species at the tips of the tree.

This might be done in various ways. For instance, we might graph the values of a quantitative trait adjacent

to the tip labels using a bar or box, or we could plot the presence or absence of different lineages from a habitat type next to the tips of the tree. Numerous such approaches have been developed and implemented in the *phytools* package, and many of these are shown in my recent book (Revell and Harmon, 2022).

Here, I'll illustrate just one such method in which a color gradient is used to visualize trait values for a set of quantitative characters at the tips of the tree. The *phytools* implementation of this plotting method is called `phylo.heatmap`.

For this example, we'll use a phylogenetic tree and morphological trait dataset for *Anolis* lizards from Mahler et al. (2010). To load these data, the reader should run the following.

```
data(anoletree)
data(anole.data)
```

Our trait data object, `anole.data`, is a data frame with six trait columns for various phylogenetic traits. We could visualize these data directly; however, the effect of overall size would tend to obscure any interesting patterns of residual variation and covariation in body shape among the species in our tree. To control for overall size, I'll first run a phylogenetic principal components analysis (Revell, 2009) using the *phytools* function `phyl.pca`. A phylogenetic principal components analysis is similar to a regular PCA except that we account for non-independence of the information for different species in our data rotation.

```
anole.pPCA<-phyl.pca(anoletree,anole.data,mode="corr")
anole.pPCA
```

```
## Phylogenetic pca
## Standard deviations:
##      PC1      PC2      PC3      PC4      PC5      PC6
## 2.2896942 0.6674345 0.4381067 0.2997973 0.1395612 0.1026573
## Loads:
##      PC1      PC2      PC3      PC4      PC5      PC6
## SVL -0.9782776 -0.01988115 0.14487425 -0.11332244 0.0781070110 -0.051442939
## HL  -0.9736568 -0.03879982 0.13442473 -0.15596460 -0.0852979941 0.028570939
## HLL -0.9711545 0.14491400 0.02151524 0.17058611 -0.0588208480 -0.053257988
## FLL -0.9759133 -0.02087140 0.14486273 0.14149988 0.0475205990 0.062386141
## LAM -0.8299594 -0.50437051 -0.23796010 0.01194704 0.0004983465 -0.003133966
## TL  -0.8679195 0.40956428 -0.27350654 -0.05871034 0.0195584629 0.018373275
```

Our PC loadings show us the the first principal component dimension is strongly negatively correlated with all of the traits in our analysis. We could consider this the "size" axis. PC 2 is most strongly (negatively) correlated with the character "LAM", number of adhesive toepad scales called lamellae; and most positively correlated with "TL", tail length.

Let's compute the principal component scores for all of our species.

```
anole.pc_scores<-scores(anole.pPCA)
head(anole.pc_scores)
```

```
##      PC1      PC2      PC3      PC4      PC5      PC6
## ahli    -0.1747576 0.8697064 1.52379491 1.6029659 -0.23955421 -0.1626049
## allogus  0.1646585 1.4017806 1.74506491 1.5358005 -0.08089546 -0.1245855
## rubribarbus -0.4925001 1.0413268 1.45866163 1.4180850 -0.05716104 -0.1797060
## imias    -1.1608049 0.7514380 0.75822327 1.7127381 0.35013533 -0.1340347
## sagrei   -0.3486332 1.2632997 -0.05102313 0.7317455 0.37217463 -0.1484157
## bremeri  -0.9714818 0.6943196 0.11689334 0.9290039 0.43486041 -0.2304513
```

Since the *sign* of each principal component is *arbitrary* (principal components are vectors), let's simply flip the sign of PC 1 so that it switches from *negative* size to simply size.

```
anole.pc_scores[,1]<--anole.pc_scores[,1]
head(anole.pc_scores)
```

##	PC1	PC2	PC3	PC4	PC5	PC6
## ahli	0.1747576	0.8697064	1.52379491	1.6029659	-0.23955421	-0.1626049
## allogus	-0.1646585	1.4017806	1.74506491	1.5358005	-0.08089546	-0.1245855
## rubribarbus	0.4925001	1.0413268	1.45866163	1.4180850	-0.05716104	-0.1797060
## imias	1.1608049	0.7514380	0.75822327	1.7127381	0.35013533	-0.1340347
## sagrei	0.3486332	1.2632997	-0.05102313	0.7317455	0.37217463	-0.1484157
## bremeri	0.9714818	0.6943196	0.11689334	0.9290039	0.43486041	-0.2304513

Finally, let's graph our results using the `phylo.heatmap` function. Since the variance in our different principal component dimensions are quite different from PC to PC, let's standardize them to have a constant variance using `standardize=TRUE`. Finally, we can update the default color palette of the plot using the argument `colors`. Here, as we've seen earlier, I'll use the `viridis` color palette.

```
phylo.heatmap(anoletree,anole.pc_scores,standardize=TRUE,
  fsize=c(0.4,0.7,0.7),pts=FALSE,split=c(0.6,0.4),
  colors=viridisLite::viridis(n=40,direction=-1))
```

## Relationship of *phytools* to other packages

The *phytools* package has grown to become (along with *ape*, *phangorn*, and *geiger*) among the most important core packages for phylogenetic analysis in the R environment.

*phytools* depends (that is, it imports critical functionality from) a variety of other R packages.

## Conclusions

Over the past decade or so, I have developed *phytools* into becoming an important software for phylogenetic analysis, particularly comparative methods and phylogenetic visualization.

## References

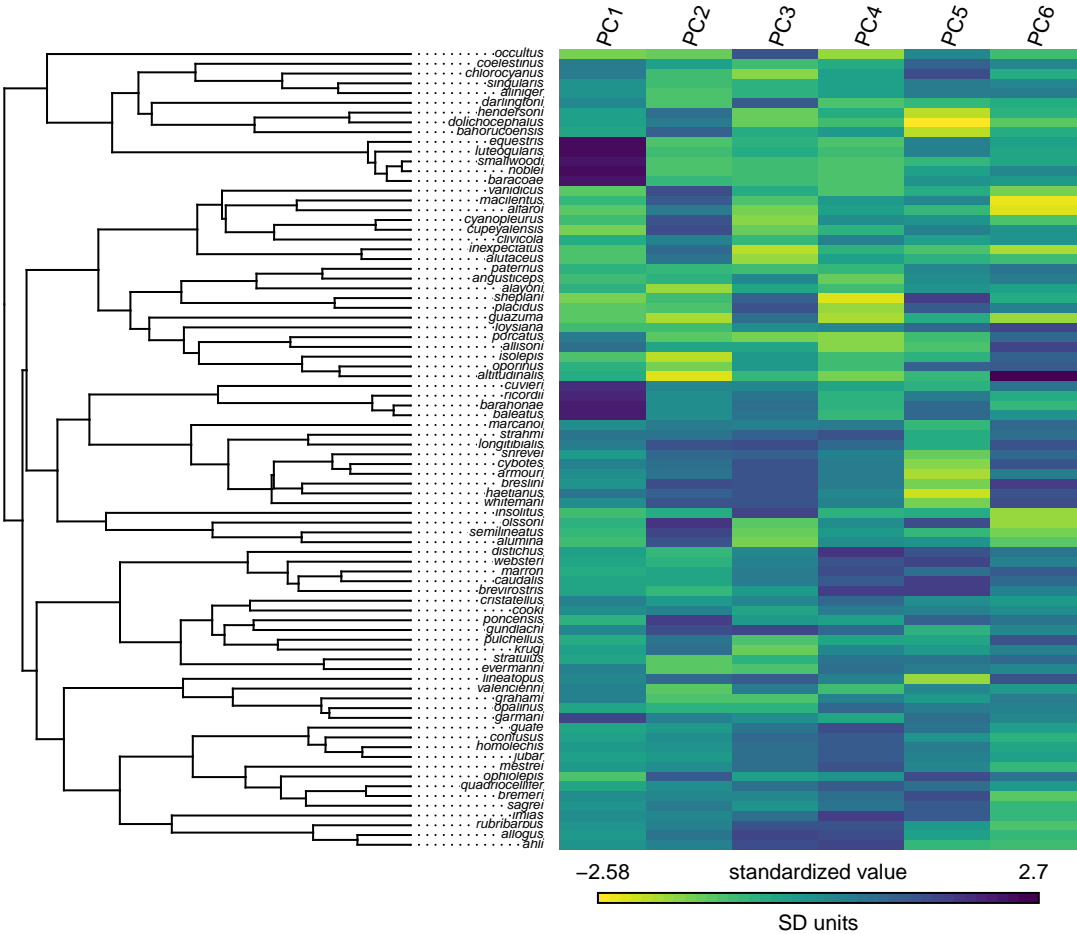


Figure 21: Phylogenetic heatmap.